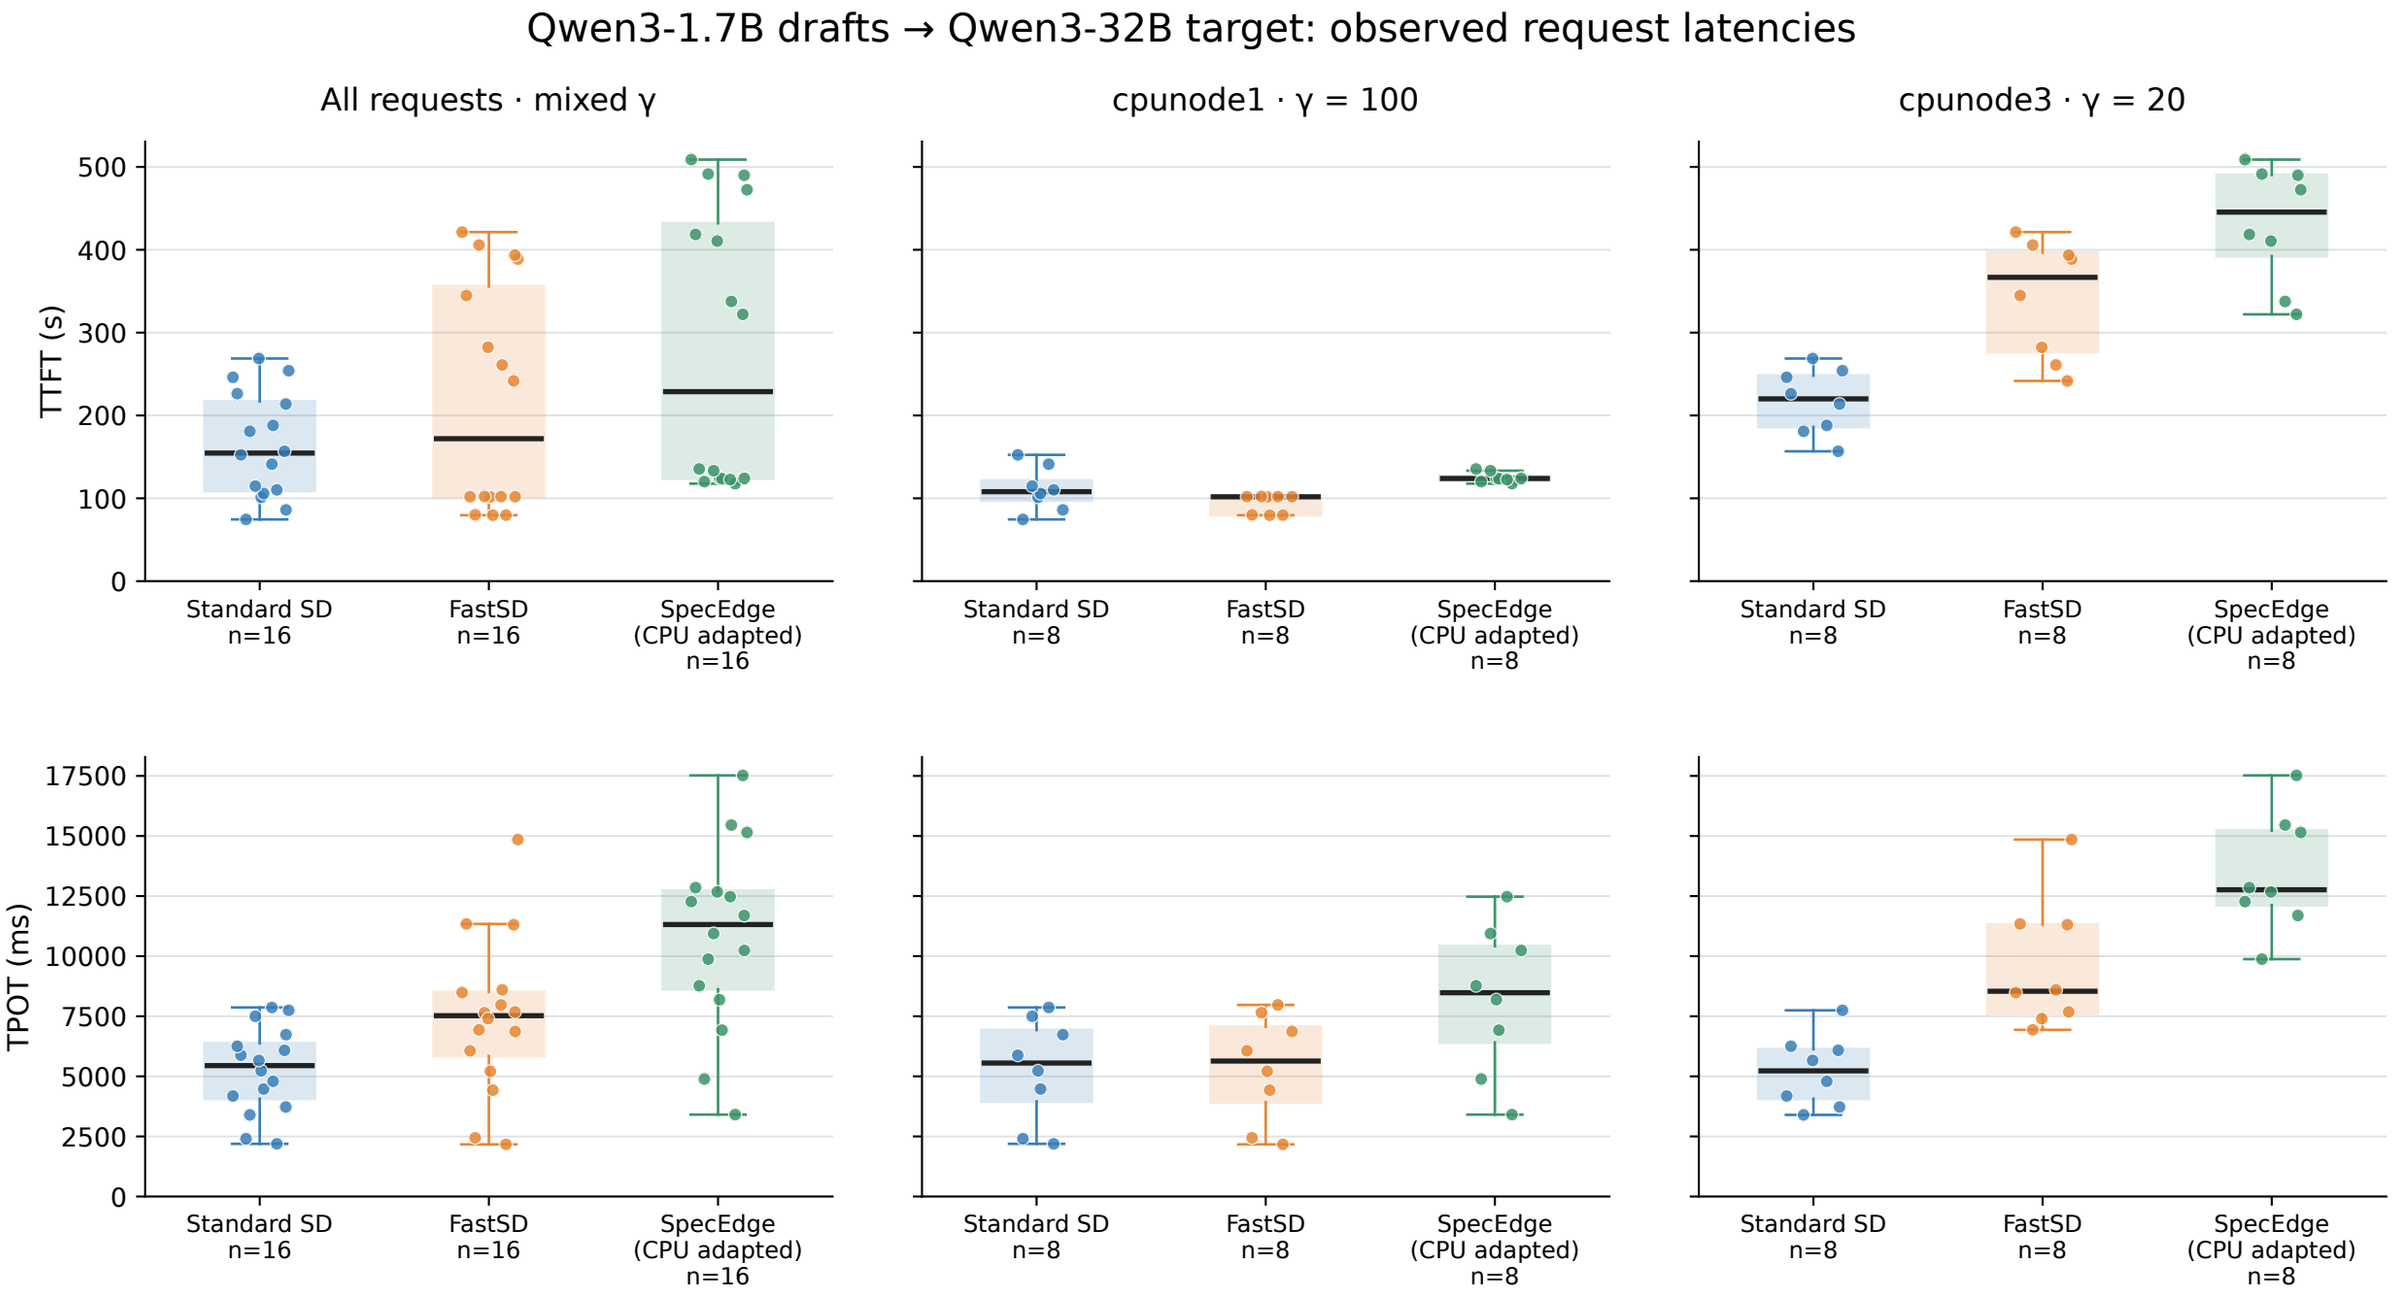

One run per method; 16 requests per method (8 per node). Dots: individual requests; boxes: IQR; black line: median; whiskers: 1.5×IQR.
Target: float16, physical GPU 3 + CPU weight offload; GPU weight budget 34 GiB. No confidence intervals are implied.
TTFT ends at the first verified output block. TPOT = (E2E − TTFT)/(output tokens − 1); single-token outputs excluded. Node, γ and prompt assignment vary together.
Host load differs between methods (see host_load.png); these observations do not isolate algorithm effects.